



Linear Regression




Colin Bennie, Caleb Hamblen, Sissi
Shen, Justin Yang



How to fit a Simple Linear Model

- Goal is to estimate:
$$y_i = B_0 + B_1 x_i + \varepsilon_i$$
- Want to minimize $(y_i - \hat{y}_i)$



Simple Linear Model Assumptions

- ϵ_i random so that the following are true:
 - $E(\epsilon_i) = 0$
 - $\text{var}(\epsilon_i) = \sigma^2$
 - $\text{Cov}(\epsilon_i, \epsilon_j) = 0$
 - $\epsilon_i \sim N(0, \sigma^2)$
- Y needs to be linear to B_0 and B_1

Results

- $E(y_i) = B_0 + B_1 x_i$: ***this is the regression function***
- $\text{var}(y_i) = \sigma^2$
- $Y_i | x_i = x \sim N(B_0 + B_1 x, \sigma^2)$
- $B_0\text{-hat} = \bar{y} - B_1\text{-hat}(\bar{x})$
- $B_1\text{-hat} = \frac{\sum(x_i - \bar{x})\sum(y_i - \bar{y})}{\sum(x_i - \bar{x})^2}$
- $Y_i\text{-hat} = B_0\text{-hat} + B_1\text{-hat}(x_i)$

How to implement in Python: Sample Data

```
In [1]: 1 import numpy as np
        2 import pandas as pd
        3 import matplotlib.pyplot as plt
        4 import statsmodels.formula.api as smf
```

```
In [2]: 1 df = pd.read_csv('Automobile_data.csv')
        2 df.head()
```

Out[2]:

	symboling	normalized-losses	make	fuel-type	aspiration	num-of-doors	body-style	drive-wheels	engine-location	wheel-base	...	engine-size	fuel-system	bore	stroke	compression-ratio	horsepower
0	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
1	3	?	alfa-romero	gas	std	two	convertible	rwd	front	88.6	...	130	mpfi	3.47	2.68	9.0	111
2	1	?	alfa-romero	gas	std	two	hatchback	rwd	front	94.5	...	152	mpfi	2.68	3.47	9.0	154
3	2	164	audi	gas	std	four	sedan	fwd	front	99.8	...	109	mpfi	3.19	3.4	10.0	102
4	2	164	audi	gas	std	four	sedan	4wd	front	99.4	...	136	mpfi	3.19	3.4	8.0	115

5 rows x 26 columns

```
In [3]: 1 df.columns
```

```
Out[3]: Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
              'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
              'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
              'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
              'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
              'highway-mpg', 'price'],
              dtype='object')
```

How to implement in Python: Model Fitting

In [8]:

```
1 import statsmodels.formula.api as smf
2
3 model = smf.ols("price ~ horsepower", data = dfnew).fit()
4 b = model.params['Intercept']
5 m = model.params['horsepower']
6 print(f'intercept={b}, slope={m}')
```

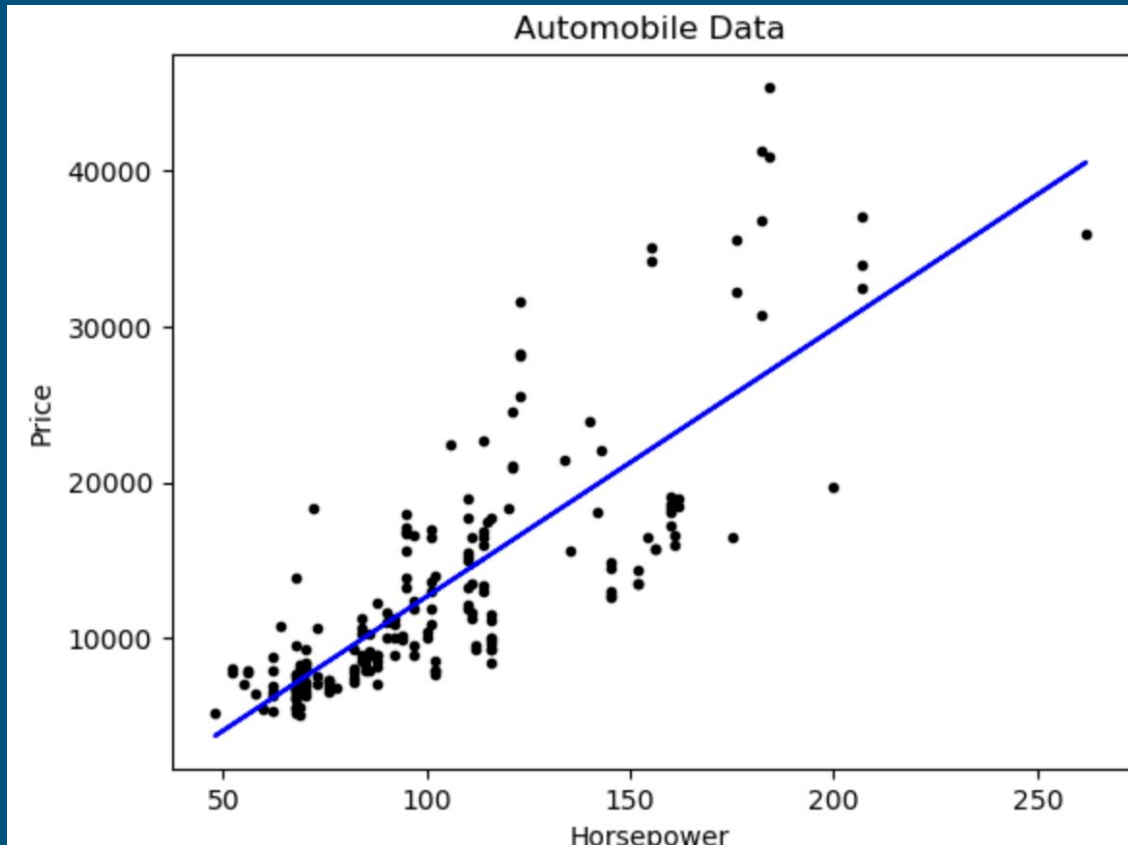
intercept=-4562.174995667492, slope=172.20625117310604

In [7]:


```
1 from sklearn import linear_model
2
3 model = linear_model.LinearRegression()
4 model.fit(dfnew[['horsepower']], dfnew['price'])
5 b = model.intercept_
6 w = model.coef_[0]
7 print(f'intercept={b}, slope={w}')
```

intercept=-4562.1749956674885, slope=172.20625117310607

How the model looks



How the model looks: residuals



Predicting using Model

In [11]:

```
1 # Sklearn
2 X = dfnew[['horsepower']]
3 y = dfnew['price']
4 model = linear_model.LinearRegression()
5 model.fit(X, y)
6
7 XX = [[1000]] # Horsepower of a Formula 1 Car
8 model.predict(XX)[0]
```

Out[11]: 167644.0761774386

In [12]:

```
1 # Statsmodels
2 model = smf.ols("price ~ horsepower", data = dfnew).fit()
3
4 XX = np.array([1000]) # Horsepower of a Formula 1 Car
5 prediction = model.params[0] + model.params[1] * XX
6 prediction[0]
```

Out[12]: 167644.07617743855

Model Diagnosis

```
In [14]: 1 model.summary()
```

Out[14]: OLS Regression Results

Dep. Variable:	price	R-squared:	0.657
Model:	OLS	Adj. R-squared:	0.655
Method:	Least Squares	F-statistic:	377.3
Date:	Fri, 29 Sep 2023	Prob (F-statistic):	1.19e-47
Time:	16:40:42	Log-Likelihood:	-1963.3
No. Observations:	199	AIC:	3931.
Df Residuals:	197	BIC:	3937.
Df Model:	1		
Covariance Type:	nonrobust		

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-4562.1750	974.995	-4.679	0.000	-6484.943	-2639.407
horsepower	172.2063	8.866	19.424	0.000	154.722	189.690
Omnibus:	38.494	Durbin-Watson:	0.749			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	64.496			
Skew:	1.013	Prob(JB):	9.89e-15			
Kurtosis:	4.916	Cond. No.	323.			

```
In [15]:
```


```
1 import statsmodels.api as sm
2
3 aov_table = sm.stats.anova_lm(model, typ=1)
4 aov_table
```

Out[15]:

	df	sum_sq	mean_sq	F	PR(>F)
horsepower	1.0	8.280790e+09	8.280790e+09	377.283543	1.189128e-47
Residual	197.0	4.323845e+09	2.194845e+07	NaN	NaN

How to fit a Multiple Linear Regression

- Uses two or more independent variables
- Fitting a plane instead of a line
- $Y = b_0 + b_1X_1 + b_2X_2 + \dots + b_nX_n + \epsilon$
- Still want to minimize $(Y_i - \hat{Y}_i)$



Multiple Linear Model Assumptions

- **Linearity:** The relationship between the dependent and independent variables is linear.
- **Independence and Normality of Errors:** The error terms (ε) are independent of each other and follow a normal distribution.
- **Homoscedasticity:** The variance of the errors is constant across all levels of independent variables.
- **No Multicollinearity:** The independent variables are not highly correlated with each other.

Results

- Regression Function: $E(Y_i) = b_0 + b_1X_{i1} + b_2X_{i2} + \dots + b_nX_{in}$
- Coefficient Estimates:
 - $B_0\text{-hat} = y\text{-bar} - B_1\text{-hat}(x\text{-bar}_1) - B_2\text{-hat}(x\text{-bar}_2) - \dots - B_n\text{-hat}(x\text{-bar}_n)$
 - $B_n\text{-hat} = \frac{\sum(x_{i_n} - x\text{-bar}_n)\sum(y_i - y\text{-bar})}{\sum(x_{i_n} - x\text{-bar}_n)^2}$
 - $B_1\text{-hat} = \frac{\sum(x_{i1} - x\text{-bar}_1)\sum(y_i - y\text{-bar})}{\sum(x_{i1} - x\text{-bar}_1)^2}$
- Fitted Values using estimated coefficients:
 - $Y_i\text{-hat} = B_0\text{-hat} + B_1\text{-hat}X_{1i} + B_2\text{-hat}X_{2i} + \dots + B_n\text{-hat}X_{ni}$

Multiple Linear Regression: Python

```
1 # Load in data for regression analysis
2 mlr_example = pd.read_csv('Automobile_data.csv')
3
4 # Identify target variable and potential predictors
5 print(mlr_example.columns)
```

✓ 0.0s

Python

```
Index(['symboling', 'normalized-losses', 'make', 'fuel-type', 'aspiration',
      'num-of-doors', 'body-style', 'drive-wheels', 'engine-location',
      'wheel-base', 'length', 'width', 'height', 'curb-weight', 'engine-type',
      'num-of-cylinders', 'engine-size', 'fuel-system', 'bore', 'stroke',
      'compression-ratio', 'horsepower', 'peak-rpm', 'city-mpg',
      'highway-mpg', 'price'],
      dtype='object')
```

```
1 # Create regression model using smf.ols
2 model = smf.ols('price ~ engine_size + horsepower + highway_mpg + fuel_type', data=model_data).fit()
3 model.summary()
```

✓ 0.0s

Python

MLR: Model Diagnosis

- R^2 and Adj- R^2
- Predictors all have their own coefficients and t test results

OLS Regression Results						
Dep. Variable:	price		R-squared:	0.812		
Model:	OLS		Adj. R-squared:	0.808		
Method:	Least Squares		F-statistic:	209.7		
Date:	Sat, 30 Sep 2023		Prob (F-statistic):	2.84e-69		
Time:	15:26:07		Log-Likelihood:	-1903.4		
No. Observations:	199		AIC:	3817.		
Df Residuals:	194		BIC:	3833.		
Df Model:	4					
Covariance Type:	nonrobust					
	coef	std err	t	P> t	[0.025	0.975]
Intercept	3071.2513	3175.980	0.967	0.335	-3192.632	9335.134
fuel_type[T.gas]	-3882.1745	910.199	-4.265	0.000	-5677.331	-2087.018
engine_size	102.6287	11.343	9.048	0.000	80.258	124.999
horsepower	58.0126	14.923	3.888	0.000	28.581	87.444
highway_mpg	-174.3590	61.706	-2.826	0.005	-296.059	-52.659
Omnibus:	11.078	Durbin-Watson:	0.841			
Prob(Omnibus):	0.004	Jarque-Bera (JB):	20.221			
Skew:	0.243	Prob(JB):	4.07e-05			
Kurtosis:	4.484	Cond. No.	2.25e+03			

MLR: Model Diagnosis

- Sequential vs. Partial F Tests

```
1 # Sequential F Test on regression model
2
3 sequential_aov_table = sm.stats.anova_lm(model, typ=1)
4 sequential_aov_table
```

✓ 0.0s

	df	sum_sq	mean_sq	F	PR(>F)
fuel_type	1.0	1.496960e+08	1.496960e+08	12.267655	5.720085e-04
engine_size	1.0	9.503851e+09	9.503851e+09	778.844709	7.616789e-70
horsepower	1.0	4.863750e+08	4.863750e+08	39.858640	1.818820e-09
highway_mpg	1.0	9.742880e+07	9.742880e+07	7.984332	5.212118e-03
Residual	194.0	2.367285e+09	1.220250e+07	NaN	NaN

Reduced Model: *price ~ fuel_type + engine_size*

Full Model: *price ~ fuel_type + engine_size + horsepower*

```
1 # Partial F Test on regression model
2
3 partial_aov_table = sm.stats.anova_lm(model, typ=2)
4 partial_aov_table
```

[43] ✓ 0.0s

	sum_sq	df	F	PR(>F)
fuel_type	2.219861e+08	1.0	18.191857	3.118291e-05
engine_size	9.989917e+08	1.0	81.867806	1.530876e-16
horsepower	1.844180e+08	1.0	15.113133	1.389653e-04
highway_mpg	9.742880e+07	1.0	7.984332	5.212118e-03
Residual	2.367285e+09	194.0	NaN	NaN

Reduced Model: *price ~ fuel_type + engine_size + highway_mpg*

Full Model: *price ~ fuel_type + engine_size + highway_mpg + horsepower*

MLR: Multicollinearity and VIF Check

```
1 # additional imports needed to check for multicollinearity
2 from statsmodels.stats.outliers_influence import variance_inflation_factor
3 from patsy import dmatrices
4
5 # Calculate the VIF for each predictor variable included in model
6 vif_y, vif_x = dmatrices('price ~ engine_size + horsepower + highway_mpg + fuel_type',
7 | | | | | data=model_data, return_type='dataframe')
8
9 vif = pd.DataFrame()
10 vif['VIF Factor'] = [variance_inflation_factor(vif_x.values, i) for i in range(vif_x.shape[1])]
11 vif['Features'] = vif_x.columns
12 print(vif)
```

✓ 0.0s

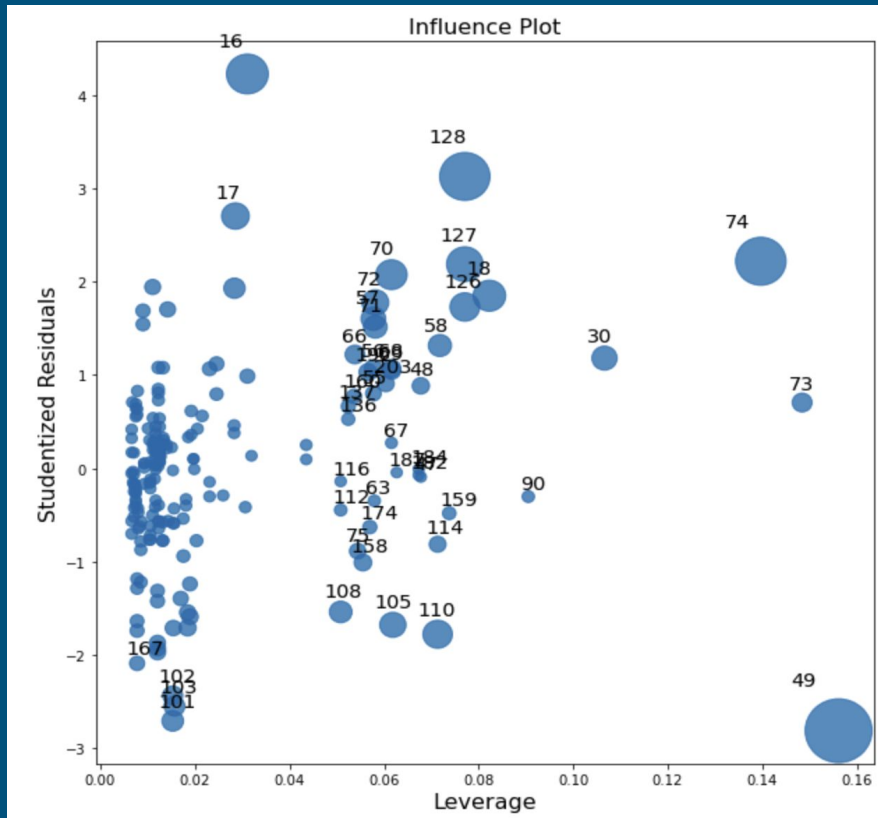
	VIF Factor	Features
0	164.497754	Intercept
1	1.221390	fuel_type[T.gas]
2	3.639274	engine_size
3	5.095827	horsepower
4	2.898497	highway_mpg

Link to Github with Code

https://github.com/cwbennie/comms_code_demo23

MLR: Influential Points

- Leverage
- Standardized Residuals
- Cook's Distance



How the model looks: Code to plot

```
In [9]: 1 import matplotlib.pyplot as plt
        2
        3 X = dfnew[['horsepower']]
        4 y = dfnew['price']
        5 model = linear_model.LinearRegression()
        6 model.fit(X, y)
        7 b = model.intercept_
        8 w = model.coef_[0]
        9
       10 x = dfnew.horsepower # we need a 1D array for plotting (and a 2D array for .fit() above)
       11 plt.plot(x, y, '.', color='black', label='data')
       12 plt.title('Automobile Data')
       13 plt.xlabel('Horsepower')
       14 plt.ylabel('Price')
       15
       16 y_hat = model.predict(X) # equivalent to  $y_{\text{hat}} = w * X[:, 0] + b$ 
       17 plt.plot(x, y_hat, color='red',
       18          label=f' $\hat{y} = \text{round}(b, 2) + (\text{round}(w, 2))x$ ')
       19
       20 plt.show()
```